

# Yongqin Wang

+1 (425)-445-5266 | [yongqin@usc.edu](mailto:yongqin@usc.edu) | Pasadena CA. 91103 | [Google Scholar](#)

## EDUCATION

---

### University of Southern California

*Ph.D. Candidate in Computer Engineering*

Los Angeles, CA

*August, 2019 – current*

### University of Southern California

*Master of Science in Computer Engineering*

Los Angeles, CA

*May, 2024*

### University of Washington

*Bachelor of Science in Electrical Engineering*

Seattle, WA

*September, 2015 – June, 2019*

## RESEARCH INTERESTS

---

- Secure Multi-party Computing
- Oblivious RAM
- Trusted Execution Environment

## INTERNSHIP EXPERIENCE

---

### Research Intern

*Meta AI*

Pasadena, CA

*May, 2021 – December 2021*

Studies major Transformer-based model inference runtime overheads and potential optimizations when Multi-party computing (MPC) is implemented. Results are published in ISPASS 2022.

## PUBLICATIONS

---

### **PIGEON: A Framework for Private Inference of Neural Networks.**

**[ under review ]**

**Authors:** Christopher Harth-Kitzerow, **Yongqin Wang**, Rachit Rajat, Murali Annavaram

This paper introduces an efficient framework for privacy-preserving MPC ML. It utilizes a novel ABG programming model that switches between Arithmetic vectorization, Bitslicing, and GPU offloading to optimize performance for different neural network layers. PIGEON significantly improves ReLU throughput by two orders of magnitude, reduces GPU memory usage, and supports larger batch sizes compared to state-of-the-art methods. It also achieves up to 70% saturation for 25Gbps networks, offering a scalable, modular, and protocol-agnostic framework for efficient private inference.

### **CompactTag: Minimizing Computation Overheads in Actively-Secure MPC for DNN.**

**[ under review ]**

**Authors:** **Yongqin Wang**, Pratik Sarkar, Nishat Koti, Arpita Patra, Murali Annavaram

MPC for active malicious adversaries uses MACs to check the final results. MAC-related computation increases runtime by around 30%. In this paper, I propose a compression technique that eliminates most of the MAC-related computations, providing up to 1.41x runtime speedups.

### **High-Throughput Secure MPC with an Honest Majority in Various Network Settings.** **[ PETS 2025 ]**

**Authors:** Christopher Harth-Kitzerow, Ajith Suresh, **Yongqin Wang**, Hossein Yalme, Georg Carle, Murali Annavaram

In this work, we present novel protocols over rings for semi-honest secure three-party computation (3PC) and malicious four-party computation (4PC) with one corruption. Our protocols address these issues by tolerating multiple arbitrarily weak network links between parties without any substantial decrease in performance. Additionally, they significantly reduce computational complexity by requiring up to half the number of basic instructions per gate compared to related work. These improvements lead to up to twice the throughput of

state-of-the-art protocols in homogeneous network settings and even larger performance improvements in heterogeneous settings.

**MPC-Pipe: An Efficient Pipeline Scheme for Semi-honest MPC Machine Learning.** [ ASPLOS 2024 ]

**Authors:** *Yongqin Wang*, Rachit Rajat, Murali Annavaram

MPC-Pipe is a novel and efficient MPC framework that utilizes communication and computation overlaps to reduce ML model inference runtime latency. There are three major pipeline schemes introduced in the paper: 1) inter-linear pipeline, 2) inner-layer pipeline, and 3) inter-batch pipeline. MPC-Pipe achieves 33% throughput improvement and 13% latency improvement for the state-of-art DNNs.

**LAORAM: A Look Ahead ORAM Architecture for Training Large Embedding Tables.** [ ISCA 2023 ]

**Authors:** *Yongqin Wang*, Rachit Rajat, Murali Annavaram

LAORAM proposes a private training method for CPU-based large embedding tables using ORAM. LAORAM proposes an aggressive superblock formation mechanism that significantly reduces the number of access to the CPU-based large embedding tables and uses a fat-tree organization to mitigate contentions over the stash in the GPU client. Those schemes combined can reduce the ORAM access latency by up to 5.4x.

**PageORAM: An Efficient DRAM Page Aware ORAM Strategy.** [ MICRO 2022 ]

**Authors:** Rachit Rajat, *Yongqin Wang*, Murali Annavaram

In this work, we introduce a new method to read paths in PathORAM to reduce stash management costs by fetching additional data nodes in the same subtree in the ORAM tree organizations to have more opportunities for data block evictions, reducing the contention to the stash space.

**Characterization of MPC-based Private Inferences for Transformer-based Models.** [ ISPASS 2022 ]

**Authors:** *Yongqin Wang*, Edward Suh, Wenjie Xiong, Benjamin Lefaudeaux, Brian Knott, Murali Annavaram, Hsien-Hsin Lee

In this work, we provide an in-depth character study of the performance overhead for running the now popular Transformer models with secure multi-party computing (MPC). Three unique challenges are identified: 1) significant Softmax runtime, 2) significant embedding table lookups, and 3) fixed point numerical stability issue.

**Byzantine-Robust and Privacy-Preserving Framework for FedML.** [ ICLR Workshop 2021 ]

**Authors:** Hanieh Hashemi, *Yongqin Wang*, Murali Annavaram

This is a federated machine learning framework that provides security against a subset of malicious clients that may send inaccurate gradient data to undermine model accuracy and convergence, and information-theoretic data privacy clients cannot access other clients' gradients. Untrusted servers only can access encoded gradients. We provide a rigorous analysis to guarantee bounds of information leakage are infinitesimally small.

**DarKnight: A Data Privacy Scheme for Training and Inference of Deep Neural Networks.** [ MICRO 2021 ]

**Authors:** Hanieh Hashemi, *Yongqin Wang*, Murali Annavaram

This work provides a cloud machine learning training & inference computation scheme that achieves input image privacy and performance improvements. We use coded computing and trusted execution environments to achieve input image privacy and GPU to achieve performance gains.

**Origami inference: Private inference using hardware enclaves.** [ IEEE CLOUD 2021 ]

**Authors:** Krishna Giri Narra, Zhifeng Lin, *Yongqin Wang*, Keshav Balasubramanian, Murali Annavaram

This work provides a cloud machine learning inference computation mode, which distributes layers in machine learning model inference to different computation units to achieve performance and input image privacy. Layers whose inputs are highly correlated with original images are computed on trusted hardware to achieve privacy and layers whose inputs have low correlation with original images are computed on GPU to improve performance. Our model can achieve considerable performance improvements on our baseline.

## ACADEMIC SERVICE

---

- Artifact evaluation committee member for ASPLOS 2025
- Artifact evaluation committee member for ISCA 2024
- Student volunteer for ISCA 2023
- Student volunteer for HPCA 2021